
MEMAUDIT:

An Exact Package-Oracle Evaluation Protocol for Budgeted Long-Term LLM Memory Writing

Anonymous Author(s)

Affiliation

Address

email

Abstract

Long-term LLM agents must decide what to write into persistent memory before future queries are known, yet existing evaluations measure only final QA accuracy, entangling memory writing with retrieval and reasoning. We introduce MEMAUDIT, a package-oracle evaluation protocol that turns memory writing into a finite, auditable optimization problem with a certified ground-truth optimum under a fixed storage budget. Each evaluation package specifies an experience stream, candidate memory representations, and future-query requirements, enabling exact measurement of how well a memory writer preserves task-relevant information. We instantiate MEMAUDIT with a semantic coverage objective under storage and exclusivity constraints and compute exact optima via certified optimization. Across controlled and naturalistic settings, MEMAUDIT disentangles representation quality, validity preservation, and budget-aware selection effects that end-to-end QA cannot isolate. This enables, for the first time, principled and reproducible evaluation of memory-writing policies independent of downstream retrieval and reasoning.

1 Introduction

LLM agents increasingly operate across sessions: they converse with users, call tools, edit code, inspect documents, and later need to reuse what happened. Long-term agent-memory systems such as MemGPT, MemoryBank, Mem0, and A-Mem already make persistent memory a first-class system component [9, 15, 3, 13]. A fixed context window makes the naive policy of retaining everything impossible. The agent must compile an experience stream into persistent memory, choosing not only whether to store an item but also which representation should survive: raw span, extracted fact, temporal event, graph edge, summary, rule, skill, compound update, or tombstone.

This paper argues that long-term memory writing should be evaluated as finite semantic compression, not only as an architecture choice. Existing long-memory benchmarks primarily score final behavior after retrieval and answering [7, 12], while MemSim/MemDaily constructs reliable personal-memory QA pairs for downstream evaluation [14]. The load-bearing object in MEMAUDIT is a package with an exact package oracle: for a fixed budget, finite candidate memories, and fixed future-query requirements, how close is a writer to the best package-feasible store? This separates write quality from downstream retrieval and reader behavior. A system may fail because it wrote the wrong memory, because retrieval missed a good memory, or because the reader ignored evidence; final QA accuracy alone cannot identify which layer failed.

The distinction matters because most deployed memory systems are evaluated only after several other choices have intervened. A memory writer may extract a useful fact but store it in a form

MemAudit write-time evaluation




Figure 1: MEMAUDIT evaluates the written memory store before retrieval and reader reasoning. The frozen package defines the exact write-time denominator, while downstream QA is measured separately through the optional retriever/reader branch.

too expensive to keep under the target budget. A retriever may miss the relevant memory even if the writer preserved it, as in retrieval-augmented pipelines where query-time evidence selection is itself a separate component [6]. A reader may see the evidence and still answer incorrectly. These are different engineering failures. MEMAUDIT isolates the first one by freezing a finite candidate set, a storage accounting rule, and an evidence-unit objective. The resulting denominator is not a claim about all possible memories; it is a reproducible answer to the narrower question “given this package, how much of the package-feasible semantic value did the writer preserve?”

Contributions. We make four evaluation-and-dataset contributions. (i) We define a finite MEMAUDIT package for budgeted memory writing with a package-oracle ratio and a union denominator for external stores. (ii) We instantiate the known concave-over-modular submodular coverage structure for semantic memory preservation. (iii) We provide exact/certified package optima and a package-oracle reference baseline, with an independent MILP audit. (iv) We release an initial MEMAUDIT suite spanning controlled exact packages, validity-heavy stress packages, human-validated model-adjudicated natural support slices with reported inter-annotator agreement, and exported-system packages for Mem0, A-Mem, and Letta.

E&D artifact role. MEMAUDIT is designed to complement systems such as Mem0, Letta, A-Mem, MemGPT-style archival memory, and graph-based memories by scoring their exported writes under a shared denominator. The released artifact contains deterministic exact-package generators, certified solvers, cached natural package exports, adjudication summaries, exported Mem0/A-Mem/Letta memory stores, and scripts to reproduce the main tables and figures without additional API calls. API-backed construction is needed only to regenerate natural annotations or rerun external memory exports.

Relation to MemSim. MemSim/MemDaily [14] automatically constructs reliable QA pairs for evaluating personal-assistant memory through downstream answering. MEMAUDIT targets a complementary layer: it evaluates the write-time memory store itself by defining finite candidate representations, costs, evidence units, and an exact budgeted optimum. Thus MemSim asks whether an agent answers generated memory questions correctly, while MEMAUDIT asks how close a written memory store is to the best package-feasible semantic store before retrieval and reader reasoning are invoked. This also lets MEMAUDIT localize where the problem sits: candidate generation, representation choice, budget-aware selection, retrieval, or reader use.

65 2 MEMAUDIT Package

66 An LLM agent observes experiences $\mathcal{E} = (e_1, \dots, e_T)$. Each experience can generate multiple can-
 67 didate write choices, such as raw text, an atomic fact, a summary, a graph edge, a rule, a compound
 68 update, or a tombstone. Let

$$\mathcal{U} = \{(i, j) : i \in [T], j \in J_i \setminus \{\text{discard}\}\}$$

69 be the virtual ground set of non-discard experience-representation choices. Each element $u = (i, j)$
 70 has cost $c_u > 0$, and $G_i = \{(i, j) : j \in J_i \setminus \{\text{discard}\}\}$ is the group of choices for experience i . A
 71 feasible memory store satisfies one storage budget and one choice per experience:

$$\mathcal{F}_B = \left\{ X \subseteq \mathcal{U} : \sum_{u \in X} c_u \leq B, \quad |X \cap G_i| \leq 1 \forall i \right\}. \quad (1)$$

72 Thus the feasibility structure is the intersection of one knapsack constraint and one partition matroid,
 73 a standard combinatorial structure for constrained selection problems [8, 11].

74 This virtual-ground-set view makes representation choice explicit instead of treating memory as a
 75 homogeneous list of strings. For example, suppose a user first says they prefer vegetarian travel
 76 meals and later says they are pescatarian now. A future query asks what meal should be booked.
 77 The package may contain a stale vegetarian fact, a current pescatarian fact, a raw span containing
 78 both turns, a tombstone invalidating the old preference, and a compound update saying “vegetarian
 79 is superseded by pescatarian.” These candidates have different costs and cover different evidence
 80 units: current truth, invalidation, temporal order, or raw provenance. The constraint $|X \cap G_i| \leq 1$
 81 prevents the package optimum from keeping every representation of one experience and makes
 82 the denominator comparable to a writer that chooses one persistent form. Because some deployed
 83 systems keep raw+summary+fact variants, the artifact also supports a relaxed $|X \cap G_i| \leq k$ audit;
 84 on Natural-87, allowing $k = 2$ changes exact OPT by only 1.3–6.0% across the evaluated budgets
 85 (Appendix C).

86 The package also fixes storage accounting. In our experiments, the default cost is normalized word-
 87 equivalent storage over the serialized memory exposed to the writer/reader. This is a declared bench-
 88 mark cost, not a universal systems-cost model: real systems also pay embedding-index slots, graph-
 89 node overhead, JSON metadata, and retrieval compute. To check that the main diagnostics are not an
 90 artifact of whitespace word counts, Appendix C recomputes the Natural-87 package and exported-
 91 store pruning under an alternative cost $c(u) = 8 + \lceil \text{bytes}(u)/24 \rceil$, approximating fixed per-record
 92 overhead plus serialized payload. The qualitative exported-system ordering is stable under this
 93 rule. When external systems export memories outside the package candidate set, those memories
 94 are added to a union package with their own measured costs rather than forced into the original
 95 candidate taxonomy.

96 The benchmark defines semantic evidence units $\mathcal{R} = \{r_1, \dots, r_M\}$. A future query q has required
 97 evidence units $R(q) \subseteq \mathcal{R}$ and nonnegative importance weights, inducing evidence weights $w_r \geq 0$.
 98 Each candidate memory has a nonnegative coverage row $a_{ur} \in [0, 1]$. The utility of a store is

$$F(X) = \sum_{r \in \mathcal{R}} w_r h_r \left(\sum_{u \in X} a_{ur} \right), \quad (2)$$

99 where each h_r is concave, nondecreasing, and normalized. The default package objective uses
 100 $h_r(z) = \min(1, z)$, so duplicate memories of the same evidence unit have diminishing marginal
 101 value.

102 Evidence units are the semantic atoms used by the benchmark. In a synthetic package they are gen-
 103 erated from the hidden event graph; in a natural support slice they are source-backed units extracted
 104 from the support sessions and then mapped to future-query requirements. A unit can represent a
 105 fact, a temporal relation, an entity preference, a deletion, an abstention condition, or a validity-state
 106 update. The objective is positive coverage: stale-fact avoidance is represented by covering the ev-
 107 idence unit that says an older fact is no longer current, not by assigning negative utility to stale
 108 memories. This keeps the objective monotone while still allowing the benchmark to ask whether
 109 writers preserve current-truth information.

110 The concavity in Equation (2) encodes diminishing returns. If two candidates cover the same ev-
 111 idence unit, the second copy should usually help less than the first. If a query requires multiple

112 evidence units, the weights w_r distribute value across those requirements. The objective is therefore
 113 a semantic surrogate for write-time preservation, not a replacement for downstream QA. A reader
 114 can still fail with a high- F store, and a low- F store might answer an easy question by chance. The
 115 point is that F gives a deterministic, package-local target for the writer layer.

116 **Theorem 1** (Semantic coverage is monotone submodular). *Let $w_r \geq 0$, $a_{ur} \geq 0$, and let each h_r be*
 117 *concave, nondecreasing, and satisfy $h_r(0) = 0$. Then F in Equation (2) is normalized, monotone*
 118 *nondecreasing, and submodular on $2^{\mathcal{U}}$.*

119 The proof is the standard concave-over-modular diminishing-returns argument used throughout sub-
 120 modular selection [8] and is given in Appendix A. The theorem supports the semantic surrogate; it
 121 does not claim black-box LLM answer accuracy is submodular.

122 **Definition 2** (MEMAUDIT package and package ratio). *An MEMAUDIT package is a tuple*

$$\mathcal{P} = (\mathcal{U}, \mathcal{G}, c, \mathcal{R}, A, w, B).$$

123 *Its exact package optimum and package ratio are*

$$\text{OPT}_{\mathcal{P}}(B) = \max_{X \in \mathcal{F}_B(\mathcal{P})} F_{\mathcal{P}}(X), \quad \rho_{\mathcal{P}}(X) = F_{\mathcal{P}}(X)/\text{OPT}_{\mathcal{P}}(B).$$

124 **Definition 3** (Union denominator for external stores). *If an external memory system writes memories*
 125 *Y not contained in $\mathcal{U}$, we evaluate them in the finite union package $\mathcal{P}^+(Y)$ obtained by adding Y ,*
 126 *their costs, and their adjudicated coverage rows. The external-store ratio is*

$$\rho_{\text{union}}(Y) = \frac{F_{\mathcal{P}^+(Y)}(Y)}{\text{OPT}_{\mathcal{P}^+(Y)}(B)}.$$

127 *We also report an analysis-only upper-pruned bound over subsets of Y to separate extraction quality*
 128 *from budget-aware selection.*

129 The union denominator is essential for scoring real systems. Suppose Mem0 writes a memory that is
 130 not one of our package candidates. Scoring it only against the package candidate optimum would be
 131 ambiguous: the system may have created a useful representation that the package did not contain. In
 132 $\mathcal{P}^+(Y)$, the exported memory becomes a first-class candidate with an adjudicated coverage row. The
 133 numerator scores exactly what the system wrote and retained, while the denominator asks what the
 134 best budget-feasible subset could have achieved using both package candidates and system exports.
 135 This makes low scores interpretable: a low raw exported-store ratio with a high upper-pruned bound
 136 indicates that the system extracted useful content but did not select or compact it well under budget.

137 3 Exact Optima and Reference Writers

138 Exact optimization is central to MEMAUDIT. For small packages, we compute $\text{OPT}_{\mathcal{P}}(B)$ by
 139 branch-and-bound over experience-representation assignments. For the default clipped-coverage
 140 objective, we also use a MILP certificate with binary candidate variables x_u and coverage variables
 141 $y_r \leq \sum_u a_{ur} x_u$, $y_r \leq 1$, following the standard separation between exact/certified optima and
 142 approximation baselines in submodular knapsack settings [11, 4]. Greedy or learned references are
 143 never labeled as OPT.

144 The branch-and-bound solver searches over groups rather than over unconstrained candidate subsets.
 145 At each group it branches over discard or one representation choice, tracks remaining budget, and
 146 uses an optimistic fractional bound over current marginal gains to prune subtrees. The bound is
 147 admissible for the clipped coverage objective because future marginal coverage can only decrease
 148 as more evidence units are covered. This solver is sufficient for the exact-small and adjudicated
 149 natural packages we report. The MILP audit is included to reduce the risk that an implementation
 150 detail in the custom solver defines the benchmark.

151 The package-oracle reference baseline is grouped value-threshold (GVT). For each arriving expe-
 152 rience, GVT forms the budget-feasible candidates whose package-oracle marginal density exceeds
 153 a threshold, then inserts the admissible representation with largest raw marginal value. A thresh-
 154 old grid gives a conservative insertion-only constant-factor guarantee under exact marginals and
 155 small-item assumptions; the theorem and proof are in Appendix A. GVT is included as a calibra-
 156 tion baseline: because it has access to exact package marginals, it helps verify that the package and
 157 solver behave sensibly. It is not the proposed deployed writer.

Table 1: Exact-solver certification. PuLP MILP and pure-Python branch-and-bound were run on the same requested-scope audit instances; equality is objective-value equality, allowing tied optimal stores with different candidate ids.

Audit	Rows	Objective matches	Max diff
B&B vs MILP	1,200	1,200/1,200	0.0

Density-only representation choice is included as a negative control. It can be arbitrarily bad even for modular utility: a tiny candidate can have higher density while losing nearly all value. Controlled experiments below show this failure empirically.

This calibration role is useful because it makes representation-choice effects visible. A pure density rule may prefer a cheap but narrow memory; a pure value rule may spend the budget too quickly on expensive raw spans. GVT combines a density threshold with within-group value choice, which is well aligned with the one-representation-per-experience structure, while remaining secondary to the benchmark artifact itself.

Proposition 4 (Density-only can be arbitrarily bad). *For every $\eta > 0$, there is a one-experience modular instance satisfying $0 < c_u \leq B/2$ in which a max-density representation rule obtains at most η times the finite package optimum.*

Theorem 5 (Scoped package-oracle GVT guarantee). *Let F be normalized, monotone, and submodular on $\mathcal{U}$. Suppose candidates arrive group-by-group, the algorithm is insertion-only, computes exact package-oracle marginals, enforces Equation (1), and chooses the maximum raw marginal among threshold-admissible candidates in each arriving group. If every non-discard representation has cost $0 < c_u \leq B/2$ and the threshold grid contains*

$$\lambda \in \left[(1 - \varepsilon) \frac{\text{OPT}}{2B}, \frac{\text{OPT}}{2B} \right],$$

then the best grid solution satisfies $F(X) \geq (1 - \varepsilon)\text{OPT}/4$. This is a calibration result for the package-oracle baseline; deployed writers without exact marginals are evaluated empirically.

4 Controlled Exact-OPT Packages

MEMAUDIT-Small generates compact hidden event graphs with raw, fact, summary, tombstone, and compound-update candidates. The exact 500-seed budget sweep maps fractional budgets 0.01, 0.02, 0.05, 0.10, 0.20 to absolute budgets $B = 1, 2, 4, 8, 16$; we discuss the non-degenerate budgets $B = 2, 4, 8, 16$.

Each package instance contains enough redundancy to make representation choice nontrivial. Raw spans are broad but expensive; atomic facts are compact but may miss temporal or validity relations; summaries can cover several units at intermediate cost; tombstones and compound updates encode current-state transitions that are not recoverable from a stale fact alone. Baselines are constrained to the same package and budget. Recency raw keeps recent raw memories until the budget is exhausted. Fact-only and summary-only restrict the candidate set to one representation family. Density-only chooses the candidate with highest current marginal value per unit cost. No-tombstone OPT removes validity candidates and then resolves the exact problem, measuring the frontier loss from omitting that representation class.

Figure 2 shows that the exact package denominator separates memory-writing policies. MEMAUDIT-GVT covers all invalidation units selected by OPT at the reported budgets (mean invalidation coverage 2.00 for each of $B = 2, 4, 8, 16$). Density-only covers no invalidation units at $B = 2, 4, 8$ and reaches only 0.264–0.361 of OPT there. Tombstone/compound-update candidates expand the feasible frontier: exact no-tombstone OPT reaches 0.523, 0.654, 0.689, 0.815 of full OPT across $B = 2, 4, 8, 16$.

The exact sweep is the cleanest evidence that MEMAUDIT measures the intended layer. The denominator is certified, all methods face the same candidate set and budget, and the gaps are not downstream retrieval artifacts. The strongest negative result is not that one heuristic loses on one seed, but that a plausible local rule systematically discards the validity-state units that exact OPT




Figure 2: Exact-small budget sweep. MEMAUDIT-GVT remains near certified package OPT, while density-only, recency raw, and single-representation baselines degrade under tight storage. Bands are bootstrap 95% intervals from the canonical 500-seed run.




Figure 3: Validity-heavy stress visualization for the headline exact distributions: base, update-chain, and temporal-interval.

200 preserves. The no-tombstone OPT curve is also important because it is not an algorithm comparison: even the best package-feasible store is worse when the candidate generator lacks current-state
 201
 202 representations. This distinguishes candidate-generation bottlenecks from selection bottlenecks.

203 5 Validity-State Stress Packages

204 Memory systems often fail by preserving stale facts. We encode current truth, supersession, deletion,
 205 and abstention as positive validity-state evidence units rather than negative utility. A tombstone
 206 therefore covers an invalidation unit such as “the old address is no longer current,” while stale
 207 retrieval and stale answers remain separate diagnostic metrics.

208 This design choice is what lets validity fit into a monotone objective. A memory store is rewarded
 209 for preserving the evidence needed to answer “what is true now?” or “should the agent abstain
 210 because the previous answer is invalid?” It is not directly penalized for retaining old memories. A
 211 stale raw span may still be useful if paired with a later tombstone or compound update, because the
 212 reader can infer the update sequence. Without such validity-state candidates, a writer can spend the
 213 budget on high-density facts and still fail on future queries that ask for current truth.

214 The stress suite sharpens the validity-memory claim. At $B = 6$, no-tombstone OPT falls to 0.513 on
 215 update_chain and 0.700 on temporal_interval, so validity-state candidates expand the

package-feasible frontier in controlled update-heavy instances. Other stressors are useful diagnostics, but the main paper restricts claims to the distributions in Figure 3.

The three headline distributions cover different validity pressures. The base distribution has ordinary preferences and events with occasional updates. The `update_chain` distribution creates repeated supersession, where the useful memory is often the latest current-state representation plus evidence that older states are obsolete. The `temporal_interval` distribution stresses interval boundaries and recency of truth, so compact temporal updates matter more than raw recency alone. The full stress suite also includes density traps, scope shifts, redundancy-heavy packages, and hard abstention cases, but those are deferred because the central main-text claim is narrower: validity-state representation choices can expand the package-feasible semantic frontier.

6 Natural and Exported-System Packages

Natural packages demonstrate that the same finite-package schema can be populated from realistic conversation support slices. We built a 200-example LongMemEval-style support-slice package with Gemini Flash-Lite through OpenRouter. Candidate generation sees answer-support sessions; query and gold-answer fields are used only in the separate benchmark-labeling pass that maps evidence units to query requirements. The package is support-sliced, not full-haystack LongMemEval, and the labels are source-backed model-adjudicated labels with explicit rejection and audit paths.

The natural construction route is intentionally auditable. The pipeline separates three artifacts: candidate memories, evidence units, and query requirements. Candidate memories are generated from support sessions. Evidence units are grounded in the same source text. Query requirements are assigned in a separate labeling pass. This separation reduces direct answer leakage into the writer while still allowing exact package scoring after the artifact is frozen. The resulting Natural-87 subset is deliberately modest: it is small, support-sliced, excludes 13 ambiguous examples from the stricter 100-example adjudication pass, and still depends on model-adjudicated evidence-unit construction. We therefore use it as a first audited natural package and exported-system diagnostic, not as a final full-haystack long-memory benchmark.

Table 2: Natural package construction reliability. The package format exposes which examples are schema-valid, adjudicated for scoring, or flagged for ambiguity instead of silently mixing them into one benchmark pool.

Artifact	Scope	Resolved/agreement	Disputed	Role
Natural-200 primary	200 ex.	200 schema-valid	–	scalable package construction
Adjudicated subset	100 ex.	87 accepted/corrected	13	main natural scoring set
Flash-Lite spot-check	30 ex.	29 accepted/corrected	1	independent consistency check
Human-edited seed set	100 ex.	100 schema-valid	–	prompt/schema development

Table 2 is intentionally conservative. The primary Natural-200 export used 681 API calls, 1,978,327 tokens, and about \$0.907. An initial secondary audit surfaced ambiguous required-unit cases, motivating a stricter Gemini-Flash adjudication pass on 100 examples; it accepted or corrected 87 examples and rejected 13 for unresolved ambiguity. The independent 30-example spot-check accepted/corrected 29 examples. Rejections mostly involved underspecified query requirements, support spans that licensed multiple plausible evidence decompositions, or candidate memories whose coverage depended on borderline summary entailment. This is an advantage of the package format: ambiguous evidence requirements become measurable dataset objects rather than hidden benchmark noise.

The coverage audit in Table 3 is part of the artifact, not a side note. On the 87-example adjudicated subset, two human annotators labeled all nonzero coverage cells plus a stratified 20% sample of model-zero cells. Human agreement is substantial, Gemini agrees closely with the human-consensus subset, and re-scoring package writers with human labels changes ratios by only 0.035 on average while preserving rankings. All natural-package rows have bootstrap intervals in the artifact; for example, Estimated-GVT at $B = 100$ is 0.829 [0.761, 0.887] under Gemini labels and 0.828 [0.755, 0.893] under human-consensus labels. The Sonnet rows are useful as a stress test: rankings remain stable, but binary labels are not interchangeable with the human-validated Gemini labels.

Table 3: Coverage-label audit summary on the 87-example adjudicated subset. Cell-level rows use the released human audit frame; package-pass rows compare model package adjudications on the same cells. Ratio MAD is the mean absolute ratio change at $B = 30, 60, 100$ where a package-level re-score is defined.

Audit pair	Scope	Cohen’s κ	Raw agree.	Ratio MAD	Ranking
Human-human	1,975 cells	0.831	0.918	–	–
Human-Gemini	1,813 agreed cells	0.863	0.934	0.035	preserved
Sonnet-Gemini	package pass	0.649	0.845	0.079/0.073 0.070	preserved, $\rho=1.0$
Sonnet-human	package pass	0.588	0.811	–	preserved
Sonnet-Gemini	blind cell pass	0.531	0.763	–	–
Sonnet-human	blind cell pass	0.484	0.739	–	–




Figure 4: Actual-system diagnostic on the 87-example adjudicated subset. Left: package or union ratio across budgets for package-oracle MEMAUDIT-GVT, non-oracle Estimated-GVT, and exported Mem0/Letta/A-Mem stores. Right: at $B = 100$, upper-pruned rows isolate extraction potential, while selected/pruned rows show the budget-aware selection actually achieved by each exported store.

259 The next question is whether the same package machinery can score memory systems that write in
 260 very different forms. The union denominator makes this possible: exported memories are added as
 261 candidates with measured costs and adjudicated coverage rows, so the score can distinguish extrac-
 262 tion quality from budget-aware selection. These exported-system scores are post-hoc budget-pruned
 263 diagnostics over written memories, not claims that Mem0, Letta, or A-Mem were natively deployed
 264 with our package budget or objective.

265 Figure 4 and Table 4 show that MEMAUDIT can diagnose heterogeneous memory systems without
 266 requiring a common internal representation. The non-oracle Estimated-GVT line is the more realis-
 267 tic package writer: it uses estimated values rather than hidden package-oracle marginals and reaches
 268 0.53, 0.68, 0.83 of package OPT across $B = 30, 60, 100$. Letta’s core+archival exports achieve the
 269 strongest salience-pruned union ratios on the adjudicated subset. Mem0’s high upper-pruned bound
 270 shows that useful evidence is often extracted, while the gap to salience pruning isolates budget-aware
 271 selection as the bottleneck. A-Mem preserves evidence in rich notes, but their serialized size makes
 272 them poorly matched to strict small-budget stores unless paired with compaction or admission con-
 273 trol.

274 The denominator comparison makes the union-package point concrete. Mem0’s upper-pruned pack-
 275 age ratio exceeds one because its exports add coverage not available to the original package candi-
 276 dates; the union denominator is the fairer finite comparison once those memories are admitted
 277 as candidate writes. Conversely, scoring only the system’s own memories would remove the opti-
 278 mum and obscure whether the bottleneck is extraction or budget-aware selection. The exactly zero
 279 A-Mem full-store row is a cost-model diagnostic rather than a coverage bug: at $B = 100$, no full
 280 serialized A-Mem note fits the budget (mean minimum full-note cost 1768 words, or 464 units under
 281 the byte-overhead rule), while its compact metadata view remains feasible.

Table 4: External-system denominator comparison at $B = 100$ on Natural-87. $F(Y)$ is the raw mean coverage value of the post-hoc pruned exported store. The package ratio uses the original package-candidate denominator; the union ratio uses the expanded denominator after adding the system’s own exported memories.

Exported diagnostic row	Raw $F(Y)$	Package ratio	Union ratio
Mem0 salience-pruned	0.655	0.489	0.427
Mem0 upper-pruned	1.443	1.041	0.886
Letta salience-pruned	1.270	0.881	0.734
Letta upper-pruned	1.322	0.918	0.765
A-Mem metadata view	0.282	0.218	0.180
A-Mem full-store view	0.000	0.000	0.000

The diagnostic also avoids a common evaluation shortcut. It would be easy to ask each memory system to answer downstream questions and report the final score. That number is useful, but it does not say whether a low score came from writing, retrieval, prompt formatting, model capability, or answer normalization. In Figure 4, the same evidence objective, costs, and budgets are used for all exported stores. This makes the failure mode more actionable: improve extraction if the upper-pruned bound is low; improve admission or compaction if the bound is high but the selected-store ratio is low; improve retrieval/reader behavior only after the written store is known to contain the necessary evidence.

7 Related Work

Memory systems and benchmarks. MemGPT, Reflexion, MemoryBank, Mem0, and A-Mem propose long-term memory architectures or policies [9, 10, 15, 3, 13]. LoCoMo, LongMemEval, and MemSim/MemDaily evaluate long-horizon memory through downstream behavior or generated QA [7, 12, 14]. MEMAUDIT is complementary: it scores exported writes under a package or union denominator before retrieval and reader errors are mixed in.

Retrieval and submodular selection. RAG, MMR, and submodular context selection choose evidence after a query is known [6, 2, 1, 5]. MEMAUDIT instead evaluates persistent write-time memory before future queries are known. Concave-over-modular objectives and knapsack submodular optimization are standard tools [8, 11, 4]; we use them to define and certify a benchmarkable semantic surrogate.

Dataset and evaluation protocols. Several long-memory datasets emphasize final answers after long histories. MEMAUDIT targets the compatible write-time layer: whether the persistent store contains needed evidence before the query-time pipeline runs. Each ratio is reproducible from a candidate table, cost vector, coverage matrix, weights, and budget.

8 Scope, Limitations, and Artifact

MEMAUDIT’s scores are package-conditional by design: the candidate table, cost model, evidence schema, and budget define the evaluated object. This supplies a write-time diagnostic rather than replacing end-to-end assistant evaluation. If a generator omits the right memory, the package optimum measures the best store within that table; if a writer extracts useful memories but fails to retain the right subset, union and upper-pruned scores expose the selection bottleneck. Costs and one-representation groups are explicit package parameters; on Natural-87, byte-overhead costs and a $k = 2$ relaxation change ratios by only a few percentage points, not qualitative rankings.

Natural-87 is deliberately small, support-sliced, and model-adjudicated, with 13 ambiguous examples rejected from the stricter adjudication pass and human audits exposing borderline entailment cases. Support slices isolate write-time memory before full-haystack expansion adds retrieval and session-selection confounds. The package-oracle reference writer uses exact coverage marginals; deployed systems need estimated utilities, learned generators, compaction, or native budget-aware admission.

References

- [1] Ashwinkumar Badanidiyuru, Baharan Mirzasoleiman, Amin Karbasi, and Andreas Krause. Streaming submodular maximization: Massive data summarization on the fly. In *Proceedings of the 20th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 671–680. ACM, 2014.
- [2] Jaime Carbonell and Jade Goldstein. The use of MMR, diversity-based reranking for reordering documents and producing summaries. In *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 335–336. ACM, 1998.
- [3] Prateek Chhikara, Dev Khant, Saket Aryan, Taranjeet Singh, and Deshraj Yadav. Mem0: Building production-ready AI agents with scalable long-term memory. *arXiv preprint arXiv:2504.19413*, 2025.
- [4] Chien-Chung Huang, Naonori Kakimura, and Yuichi Yoshida. Streaming algorithms for maximizing monotone submodular functions under a knapsack constraint. In *Approximation, Randomization, and Combinatorial Optimization. Algorithms and Techniques (APPROX/RANDOM 2017)*, volume 81 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 11:1–11:14, Dagstuhl, Germany, 2017. Schloss Dagstuhl – Leibniz-Zentrum für Informatik.
- [5] Lilly Kumari, Shengjie Wang, Tianyi Zhou, Nikhil Sarda, Anthony Rowe, and Jeff Bilmes. BumbleBee: Dynamic KV-cache streaming submodular summarization for infinite-context transformers. In *Proceedings of the First Conference on Language Modeling*, 2024.
- [6] Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33:9459–9474, 2020.
- [7] Adyasha Maharana, Dong-Ho Lee, Sergey Tulyakov, Mohit Bansal, Francesco Barbieri, and Yuwei Fang. Evaluating very long-term conversational memory of LLM agents. *arXiv preprint arXiv:2402.17753*, 2024.
- [8] George L. Nemhauser, Laurence A. Wolsey, and Marshall L. Fisher. An analysis of approximations for maximizing submodular set functions—I. *Mathematical Programming*, 14(1):265–294, 1978.
- [9] Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G. Patil, Ion Stoica, and Joseph E. Gonzalez. MemGPT: Towards LLMs as operating systems. *arXiv preprint arXiv:2310.08560*, 2023.
- [10] Noah Shinn, Federico Cassano, Edward Berman, Ashwin Gopinath, Karthik Narasimhan, and Shunyu Yao. Reflexion: Language agents with verbal reinforcement learning. *arXiv preprint arXiv:2303.11366*, 2023.
- [11] Maxim Sviridenko. A note on maximizing a submodular set function subject to a knapsack constraint. *Operations Research Letters*, 32(1):41–43, 2004.
- [12] Di Wu, Hongwei Wang, Wenhao Yu, Yuwei Zhang, Kai-Wei Chang, and Dong Yu. Long-memeval: Benchmarking chat assistants on long-term interactive memory. *arXiv preprint arXiv:2410.10813*, 2024.
- [13] Wujiang Xu, Zujie Liang, Kai Mei, Hang Gao, Juntao Tan, and Yongfeng Zhang. A-MEM: Agentic memory for LLM agents. *arXiv preprint arXiv:2502.12110*, 2025.
- [14] Zeyu Zhang, Quanyu Dai, Luyu Chen, Zeren Jiang, Rui Li, Jieming Zhu, Xu Chen, Yi Xie, Zhenhua Dong, and Ji-Rong Wen. MemSim: A bayesian simulator for evaluating memory of LLM-based personal assistants. *arXiv preprint arXiv:2409.20163*, 2024.
- [15] Wanjun Zhong, Lianghong Guo, Qiqi Gao, He Ye, and Yanlin Wang. Memorybank: Enhancing large language models with long-term memory. *arXiv preprint arXiv:2305.10250*, 2023.

368 A Deferred Proofs

369 *Proof of Theorem 1.* Normalization follows from $h_r(0) = 0$. For monotonicity, adding any element
 370 u increases each inner sum by $a_{ur} \geq 0$, and h_r is nondecreasing. For submodularity, let $S \subseteq T \subseteq \mathcal{U}$
 371 and $u \notin T$. Define $z_r(S) = \sum_{v \in S} a_{vr}$. Since $z_r(S) \leq z_r(T)$ and h_r is concave, the increment
 372 $h_r(z + a_{ur}) - h_r(z)$ is nonincreasing in z for every $a_{ur} \geq 0$. Multiplying by $w_r \geq 0$ and summing
 373 over r gives $\Delta(u \mid S) \geq \Delta(u \mid T)$. $\square$

374 *Proof of Proposition 4.* Let $B = 2$ and choose $\delta = \min\{\eta/2, 1/4\}$. One experience has candidates
 375 a, b with $c_a = \delta$, $F(\{a\}) = 2\delta$, $c_b = 1$, and $F(\{b\}) = 1$. Candidate a has density 2, candidate b has
 376 density 1, so max-density chooses a while OPT chooses b . The ratio is $2\delta \leq \eta$. $\square$

377 *Proof of Theorem 5.* Fix such a threshold. If the final cost $C(X) \geq B/2$, every accepted element
 378 had marginal density at least λ , so $F(X) \geq \lambda C(X) \geq (1 - \varepsilon)\text{OPT}/4$. Otherwise $C(X) < B/2$.
 379 Let O be optimal. By submodularity, $\text{OPT} \leq F(X) + \sum_{u \in O \setminus X} \Delta(u \mid X)$. Any missed optimal
 380 element was budget-feasible at its arrival. If it was below threshold, its marginal charge is at most
 381 λc_u , and the total below-threshold charge is at most $\lambda B \leq \text{OPT}/2$. If it was threshold-admissible,
 382 the algorithm selected another candidate from the same group with at least as large raw marginal;
 383 these charges inject into accepted arrivals and sum to at most $F(X)$. Hence $\text{OPT} \leq 2F(X) +$
 384 $\text{OPT}/2$, so $F(X) \geq \text{OPT}/4$. $\square$

385 B Additional Experimental Details

386 This appendix records diagnostics omitted from the main story to keep the core paper within the
 387 9-page main-body budget. These results are useful for reproducibility and interpretation, but the
 388 main claims rely on exact package optima and denominator-matched exported-store scoring.

389 B.1 Controlled Package Generator

390 MEMAUDIT-Small samples a compact hidden event graph, exposes an ordered stream of experi-
 391 ences, and generates finite representation choices for each experience. Candidate types include raw
 392 span, atomic fact, entity summary, tombstone, and compound update. Costs are measured in nor-
 393 malized word-equivalent units. Evidence units correspond to future-query support requirements,
 394 including positive units for validity state: a tombstone is useful when a future query requires know-
 395 ing that an older fact has been invalidated, superseded, deleted, or should trigger abstention.


396 The benchmark reports ratios only when the denominator is certified. For the 500-seed exact-small
 397 run, each budgeted package is solved exactly with the pure-Python branch-and-bound solver. A sep-
 398 arate 1,200-row audit verifies equality with the MILP formulation for the clipped-coverage objective.
 399 The no-tombstone OPT ablation removes tombstone and compound-update candidates before solv-
 400 ing the same exact problem; it is therefore a representation-frontier comparison, not an algorithmic
 401 failure mode.

402 B.2 Reader Failure Audit

403 The reader failure audit decomposes the LongMemEval-S focus diagnostic after the memory stores
 404 have been written. It is not a new denominator: it shows how often downstream errors come from
 405 missing retrieved evidence, abstention despite support, reader errors, unsupported answers, or scor-
 406 ing/parse uncertainty. This keeps the main package-ratio claim separate from retrieval and reader
 407 behavior.

408 B.3 Natural Package Construction

409 The Natural-200 package is built from support slices rather than full conversation histories. For each
 410 example, the candidate-generation stage receives source sessions and emits candidate memories with
 411 estimated costs. A separate annotation stage maps evidence units to future-query requirements and
 412 scores candidate coverage. We intentionally separate these stages so the writer does not see the
 413 query label during memory generation.



The package can reward current-state and invalidation evidence without penalizing raw old memories directly.

Figure 5: Illustrative validity-state package. A stale fact, current fact, tombstone, and compound update are distinct candidate writes with different costs and evidence coverage. The package optimum may keep at most one candidate from an experience group, so representation choice is non-trivial rather than a union of all useful strings. The objective rewards current truth, supersession, deletion, and abstention evidence so stale-fact avoidance can be scored at write time.




Figure 6: LongMemEval-S focus reader failure audit. Bars decompose retrieval/reader outcomes after write-time memory selection; the green segment marks examples with no flagged failure in this audit.

414 The annotation results should be read as model-adjudicated dataset construction with explicit audit
415 hooks. The secondary 50-example audit surfaced ambiguous required-unit cases, which is why the
416 main paper reports the stricter 87-example adjudicated subset for system diagnostics. The most
417 common error/ambiguity categories were: required units that were too broad or could be decom-
418 posed several ways; summary candidates that implied but did not explicitly state an evidence unit;
419 graph-edge memories whose relation label was plausible but too coarse; temporal or validity updates
420 where current truth depended on ordering; and model-zero cells where humans judged a paraphrase
421 sufficient. For the released human coverage audit, two project annotators familiar with the schema
422 independently label all model-positive coverage cells on the 87-example subset plus a stratified 20%
423 sample of model-zero cells. We compute unweighted Cohen's κ on the full stratified audit frame,
424 rather than reweighting to the package's natural positive/negative balance. Human-human disagree-
425 ments are concentrated in sampled model-zero cells (116/162) and in summary/graph-edge/atomic-
426 fact boundaries; human-Gemini disagreements show the same pattern (81/119 from sampled zero
427 cells), indicating that the main ambiguity is borderline entailment rather than schema failure. The
428 human-edited seed set validates schema usability and remains in the artifact as prompt/schema de-
429 velopment data.

430 A Claude Sonnet 4.5 cross-model audit gives a useful boundary on this result. Re-adjudicating
431 the package with Sonnet preserves package-writer rankings exactly at $B = 30, 60, 100$, but pro-

duces lower cell-level agreement than Gemini–human (Table 3); we therefore treat Sonnet as a prompt/model-sensitivity stress test rather than an independent replacement annotator.

B.4 External System Exports

Mem0, A-Mem, and Letta are scored through exported written memories under a finite union denominator. The system memories are not forced into the package candidate schema; instead, they are added as new union candidates with their own costs and adjudicated coverage rows. This permits fair scoring of system-specific writes while preserving the exact denominator.

The main-system diagnostic reports post-hoc budget pruning because the public systems are not natively configured as package-budgeted writers. Recency pruning is a deployable baseline, salience pruning is a stronger heuristic, and upper pruning is an analysis-only bound over the system’s own exported memories. A high upper-pruned bound with a low salience-pruned score indicates that extraction produced useful memories but budget-aware selection failed.

C Cost and Constraint Sensitivity

We audit two package-design choices on the 87-example adjudicated natural subset. First, we relax the partition constraint from one candidate per experience to two candidates per experience. Second, we replace the default word-equivalent cost with a per-record-overhead rule $c(u) = 8 + \lceil \text{bytes}(u)/24 \rceil$. All rows below are deterministic recomputations from cached package/export artifacts and make no API calls.

Table 5: Package sensitivity on Natural-87. The $k = 2$ relaxation changes exact OPT modestly, and the byte-overhead cost increases absolute OPT because short atomic memories remain cheap relative to the tested budgets.

B	OPT $k = 1$	OPT $k = 2$	$k = 2/k = 1$	byte-overhead OPT	byte/default
30	0.931	0.943	1.013	1.310	1.322
60	1.356	1.374	1.027	1.483	1.107
100	1.489	1.540	1.060	1.632	1.131

Table 6 recomputes exported-system pruning at $B = 100$ against the package-candidate denominator under both cost rules. The exact union-denominator results in the main text remain the primary system diagnostic; this table isolates whether the small-budget ranking is mostly a word-count artifact. The qualitative ordering is stable: Letta salience remains strongest among non-oracle exported selections, Mem0 salience trails, and A-Mem full-store remains infeasible at this budget. A-Mem’s full notes are too large under both accounting rules: the mean minimum full-note cost is 1767.9 words and 463.5 byte-overhead units, far above $B = 100$.

Table 6: Exported-system cost sensitivity at $B = 100$, reported as selected value divided by package-candidate OPT under each cost rule. Upper rows use hidden coverage and are analysis-only.

Cost rule	Method	Ratio to package OPT
word	Letta salience	0.899
word	Mem0 salience	0.489
word	A-Mem metadata recency	0.190
word	A-Mem full native	0.000
byte-overhead	Letta salience	0.800
byte-overhead	Mem0 salience	0.442
byte-overhead	A-Mem metadata recency	0.266
byte-overhead	A-Mem full native	0.000




Figure 7: LongMemEval-S transfer diagnostics. These are evidence-use and reader diagnostics, not exact package-oracle ratio results. Exact-match gains are not the basis of the main claim.

D LongMemEval Transfer Diagnostics

We also ran a LongMemEval-S retrieval/reader diagnostic. These runs are not package-oracle ratios because the full benchmark does not provide a finite candidate package with exact coverage rows. They are included to check whether package-style writing can reduce retrieval cost and improve evidence use in a more natural long-memory task.

The useful signal is narrower than end-to-end QA superiority. The MEMAUDIT-style store improved retrieval compactness and token-F1/evidence-use behavior relative to full raw-store dense retrieval, while exact-match changes were not the primary evidence. This is why the main paper uses LongMemEval only as a transfer diagnostic and does not claim significant answer-accuracy gains.

E Reproducibility Notes

The artifact contains deterministic exact-package runs, natural package exports, model-adjudication reports, and public-system scoring scripts. The safest reproduction path is to run the exact-small suite against certified OPT, audit a random subset with the MILP solver, re-score the 87-example adjudicated natural subset under package and union denominators, and recompute the public-system pruning diagnostics from exported Mem0, A-Mem, and Letta stores. API-backed natural package construction caches prompts, completions, token counts, model versions, and adjudication decisions. Natural package candidate generation used OpenRouter with google/gemini-3.1-flash-lite-preview, temperature 0, JSON response format, and a default 1400-token completion cap. Exported-system coverage/saliency adjudication used cached Gemini Flash/Flash-Lite calls with temperature 0 and JSON response format; the Letta coverage and saliency caps were 1800 and 4000 tokens respectively, while the Mem0 writer configuration used temperature 0 and a 700-token cap. The cross-model stress audit used OpenRouter anthropic/claude-sonnet-4.5 with temperature 0, JSON response format, cached package-level adjudications, and a separate cached 1,975-cell binary coverage pass. All cached calls include prompt hashes and token/cost metadata. The exact package-oracle denominator itself is deterministic once a package has been constructed.

NeurIPS Paper Checklist

1. **Claims.** [Yes] The abstract and introduction state that MEMAUDIT is an exact package-oracle evaluation protocol for finite memory-writing packages, not a deployed memory architecture. The scope and package-conditional nature are discussed in [Section 8](#).
2. **Limitations.** [Yes] [Section 8](#) gives a separate scope and limitations section covering package-conditional optima, model-adjudicated natural labels, support-sliced natural packages, exact-solver scale, and downstream QA diagnostics.
3. **Theory, assumptions, and proofs.** [Yes] Assumptions for the semantic coverage theorem and scoped oracle-GVT guarantee are stated in [Sections 2](#) and [3](#); complete proofs are provided in [Appendix A](#).
4. **Experimental result reproducibility.** [Yes] The artifact includes deterministic package generators, exact solvers, cached package artifacts, result summaries, and rerun commands. Minimal reproduction does not require API calls; API-backed reconstruction is documented separately.
5. **Open access to data and code.** [Yes] The anonymized submission includes code, package data, Croissant metadata, and documentation for reproducing the main exact-package and exported-system diagnostics.
6. **Experimental setting/details.** [Yes] [Sections 4](#) to [6](#) and [appendix B](#) describe the package distributions, budgets, candidate families, adjudication path, exact solvers, and exported-system scoring setup.
7. **Experiment statistical significance.** [Yes] The controlled exact sweep reports bootstrap 95% confidence intervals over the canonical 500-seed run. Natural and exported-system package rows are reported as deterministic diagnostics on the adjudicated subset.
8. **Experiments compute resources.** [Yes] [Section 8](#) and the artifact documentation state that exact-small and cached natural re-scoring run on CPU without API calls. API-backed natural annotation and external memory exports require OpenRouter-compatible API access and cached model/version metadata.
9. **Code of ethics.** [Yes] The work is an evaluation artifact for memory-writing systems. The release is anonymized for review and excludes API keys, private review notes, and local environment files.
10. **Broader impacts.** [Yes] The main risk is that memory benchmarks can incentivize storing sensitive or stale personal information. [Sections 5](#) and [8](#) discuss validity-state memories, abstention/deletion units, and auditability as mitigation-oriented design choices.
11. **Safeguards.** [N/A] The paper does not release a pretrained model or high-risk generative model. The artifact releases package schemas, generators, scoring code, and cached benchmark data.
12. **Licenses.** [Yes] The paper cites external systems and datasets used for comparison. The artifact documentation includes dependency and asset manifests; external repository checkouts are not bundled in the anonymized artifact.
13. **Assets.** [Yes] The submission includes a dataset/evaluation artifact with README, reproducibility instructions, evaluation card, artifact manifest, Croissant metadata, and documented intended/not-intended uses.
14. **Crowdsourcing and research with human subjects.** [N/A] No paid crowd workers or human-subject experiments are used. The human-edited seed package consists of fictional examples edited for schema validation and does not include participant data.
15. **IRB approvals.** [N/A] The released examples are synthetic, model-adjudicated support slices, or fictional human-edited seed examples. No identifiable participant data or human-subject intervention is released.
16. **Declaration of LLM usage.** [Yes] LLMs are used for natural support-slice package construction, adjudication, external memory-system exports, and reader diagnostics; these uses are described in [Sections 6](#) and [8](#) and [appendices B](#) and [D](#).